

# CSC 311: Introduction to Machine Learning

## Lecture 5 - Linear Models III, Neural Nets I

---

Amanjit Singh Kainth

University of Toronto, Summer 2025

# Outline

- 1 Softmax Regression
- 2 Tracking Model Performance
- 3 Limits of Linear Classification
- 4 Introducing Neural Networks
- 5 Expressivity of a Neural Network


# Softmax Regression

---

- 1 Softmax Regression
- 2 Tracking Model Performance
- 3 Limits of Linear Classification
- 4 Introducing Neural Networks
- 5 Expressivity of a Neural Network

# Multi-class Classification

Task is to predict a discrete( $> 2$ )-valued target.



# Targets in Multi-class Classification

- Targets form a discrete set  $\{1, \dots, K\}$ .
- Represent targets as **one-hot vectors** or **one-of-K encoding**:

$$\mathbf{t} = (\underbrace{0, \dots, 0, 1, 0, \dots, 0}_{\text{entry } k \text{ is } 1}) \in \mathbb{R}^K$$

# Linear Function of Inputs

Vectorized form:

$$\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b} \text{ or}$$

$$\mathbf{z} = \mathbf{W}\mathbf{x} \text{ with dummy } x_0 = 1$$

Non-vectorized form:

$$z_k = \sum_{j=1}^D w_{kj}x_j + b_k \text{ for } k = 1, 2, \dots, K$$

- $\mathbf{W}$ :  $K \times D$  matrix.
- $\mathbf{x}$ :  $D \times 1$  vector.
- $\mathbf{b}$ :  $K \times 1$  vector.
- $\mathbf{z}$ :  $K \times 1$  vector.

# Generating a Prediction

Interpret  $z_k$  as how much the model prefers the  $k$ -th prediction.

$$y_i = \begin{cases} 1, & \text{if } i = \arg \max_k z_k \\ 0, & \text{otherwise} \end{cases}$$

How does the  $K = 2$  case relate to the binary linear classifiers?



# Softmax Regression

- Soften the predictions for optimization.
- A natural activation function is the **softmax function**, a generalization of the logistic function:

$$y_k = \text{softmax}(z_1, \dots, z_K)_k = \frac{e^{z_k}}{\sum_{k'} e^{z_{k'}}}$$

- Inputs  $z_k$  are called the logits.
- Interpret outputs as probabilities.
- If  $z_k$  is much larger than the others, then  $\text{softmax}(\mathbf{z})_k \approx 1$  and it behaves like argmax.

What does the  $K = 2$  case look like?

# Cross-Entropy as Loss Function

Use cross-entropy as the loss function.

$$\mathcal{L}_{\text{CE}}(\mathbf{y}, \mathbf{t}) = - \sum_{k=1}^K t_k \log y_k = -\mathbf{t}^\top (\log \mathbf{y}),$$

where the  $\log$  is applied element-wise.

Often use a combined **softmax-cross-entropy** function.

# Gradient Descent Updates for Softmax Regression

Softmax Regression:

$$\mathbf{z} = \mathbf{W}\mathbf{x}$$

$$\mathbf{y} = \text{softmax}(\mathbf{z})$$

$$\mathcal{L}_{\text{CE}} = -\mathbf{t}^\top (\log \mathbf{y})$$

Gradient Descent Updates:

$$\frac{\partial \mathcal{L}_{\text{CE}}}{\partial \mathbf{w}_k} = \frac{\partial \mathcal{L}_{\text{CE}}}{\partial z_k} \cdot \frac{\partial z_k}{\partial \mathbf{w}_k} = (y_k - t_k) \cdot \mathbf{x}$$

$$\mathbf{w}_k \leftarrow \mathbf{w}_k - \alpha \frac{1}{N} \sum_{i=1}^N (y_k^{(i)} - t_k^{(i)}) \mathbf{x}^{(i)}$$

## Tracking Model Performance

---

- 1 Softmax Regression
- 2 Tracking Model Performance
- 3 Limits of Linear Classification
- 4 Introducing Neural Networks
- 5 Expressivity of a Neural Network

# Progress During Learning

- Track progress during learning by plotting training curves.
- Chose the training criterion (e.g. squared error, cross-entropy) partly to be easy to optimize.
- May wish to track other **metrics** to measure performance (even if we can't directly optimize them).




Figure 1. Training error (left) and test error (right) on CIFAR-10 with 20-layer and 56-layer “plain” networks. The deeper network has higher training error, and thus test error. Similar phenomena on ImageNet is presented in Fig. 4.

# Tracking Accuracy for Binary Classification

We can track **accuracy**, or fraction correctly classified.

- Equivalent to the average 0–1 loss, the **error rate**, or fraction incorrectly classified.
- Useful metric to track even if we couldn't optimize it.

Another way to break down the accuracy:

$$Acc = \frac{TP + TN}{P + N} = \frac{TP + TN}{(TP + FN) + (TN + FP)}$$

- $P$ : num positive;  $N$ : num negative;
- $TP$ : true positives;  $TN$ : true negatives
- $FP$ : false positive or a type I error
- $FN$ : false negative or a type II error

## Accuracy is Highly Sensitive to Class Imbalance

Suppose you are screening patients for a particular disease. It's known that 1% of patients have that disease.

- What is the simplest model that can achieve 99% accuracy?



# Sensitivity and Specificity

Useful metrics even under class imbalance!


$$\text{Sensitivity} = \frac{TP}{TP+FN} \text{ [True positive rate]}$$

$$\text{Specificity} = \frac{TN}{TN+FP} \text{ [True negative rate]}$$

What happens if our problem is not linearly separable?  
How do we pick a threshold for  $y = \sigma(x)$ ?


# Designing Diagnostic Tests

- A binary model to predict whether someone has a disease.
- What happens to sensitivity and specificity as you slide the threshold from left to right?




# Tradeoff between Sensitivity and Specificity

As we increase the criterion value (i.e. move from left to right), how do the sensitivity and specificity change?




# Receiver Operating Characteristic (ROC) Curve

Area under the ROC curve (AUC) can quantify if a binary classifier achieves a good tradeoff between sensitivity and specificity.



# Confusion Matrix for Multi-Class classification

- Visualizes how frequently certain classes are confused.
- $K \times K$  matrix; rows are true labels, columns are predicted labels, entries are frequencies
- What does the confusion matrix for a perfect classifier look like?




## Limits of Linear Classification

---

- 1 Softmax Regression
- 2 Tracking Model Performance
- 3 Limits of Linear Classification
- 4 Introducing Neural Networks
- 5 Expressivity of a Neural Network

# XOR is Not Linearly Separable

Some datasets are not linearly separable, e.g. **XOR**.




Visually obvious, but how can we prove this formally?



# Proof That XOR is Not Linearly Separable

Proof by Contradiction:

- Half-spaces are convex. That is, if two points lie in a half-space, the line segment connecting them also lies in the same half-space.
- Suppose that the problem is feasible.
- If the positive examples are in the positive half-space, then the green line segment must be as well.
- Similarly, the red line segment must lie in the negative half-space.
- But, the intersection can't lie in both half-spaces. Contradiction!



# Classifying XOR Using Feature Maps

Sometimes, we can overcome this limitation using **feature maps**, e.g., for **XOR**.

$\psi(\mathbf{x}) = \begin{pmatrix} x_1 \\ x_2 \\ x_1x_2 \end{pmatrix}$	$x_1$	$x_2$	$\psi_1(\mathbf{x})$	$\psi_2(\mathbf{x})$	$\psi_3(\mathbf{x})$	$t$
	0	0	0	0	0	0
	0	1	0	1	0	1
	1	0	1	0	0	1
	1	1	1	1	1	0

- This is linearly separable. (Try it!)
- Designing feature maps can be hard. Can we learn them?


# Introducing Neural Networks

---

- 1 Softmax Regression
- 2 Tracking Model Performance
- 3 Limits of Linear Classification
- 4 Introducing Neural Networks**
- 5 Expressivity of a Neural Network

# Neurons in the Brain


Neurons receive input signals and accumulate voltage. After some threshold, they will fire spiking responses.



[Pic credit: [www.moleculardevices.com](http://www.moleculardevices.com)]

# A Simpler Neuron


For neural nets, we use a much simpler model for neuron, or **unit**:



- Similar to logistic regression:  $y = \sigma(\mathbf{w}^\top \mathbf{x} + b)$
- By throwing together lots of these simple neuron-like processing units, we can do some powerful computations!

# A Feed-Forward Neural Network

- A **directed acyclic graph (DAG)**
- Units are grouped into **layers**




# Multilayer Perceptrons

- A multi-layer network consists of fully connected layers.
- In a fully connected layer, all input units are connected to all output units.
- Each hidden layer  $i$  connects  $N_{i-1}$  input units to  $N_i$  output units. Weight matrix is  $N_i \times N_{i-1}$ .
- The outputs are a function of the input units:


$$\mathbf{y} = f(\mathbf{x}) = \phi(\mathbf{W}\mathbf{x} + \mathbf{b})$$

$\phi$  is applied component-wise.






# Some Activation Functions




Identity

$$y = z$$



Rectified Linear Unit  
(ReLU)


$$y = \max(0, z)$$



Soft ReLU


$$y = \log 1 + e^z$$

# More Activation Functions




Hard Threshold

$$y = \begin{cases} 1 & \text{if } z > 0 \\ 0 & \text{if } z \leq 0 \end{cases}$$



Logistic

$$y = \frac{1}{1 + e^{-z}}$$



Hyperbolic Tangent  
(tanh)

$$y = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$

# Computation in Each Layer

Each layer computes a function.

$$\mathbf{h}^{(1)} = f^{(1)}(\mathbf{x}) = \phi(\mathbf{W}^{(1)}\mathbf{x} + \mathbf{b}^{(1)})$$

$$\mathbf{h}^{(2)} = f^{(2)}(\mathbf{h}^{(1)}) = \phi(\mathbf{W}^{(2)}\mathbf{h}^{(1)} + \mathbf{b}^{(2)})$$

$\vdots$


$$\mathbf{y} = f^{(L)}(\mathbf{h}^{(L-1)})$$

If task is regression: choose

$$\mathbf{y} = f^{(L)}(\mathbf{h}^{(L-1)}) = (\mathbf{w}^{(L)})^\top \mathbf{h}^{(L-1)} + b^{(L)}$$

If task is binary classification: choose

$$\mathbf{y} = f^{(L)}(\mathbf{h}^{(L-1)}) = \sigma((\mathbf{w}^{(L)})^\top \mathbf{h}^{(L-1)} + b^{(L)})$$




# A Composition of Functions

The network computes  
a composition of functions.


$$\mathbf{y} = f^{(L)} \circ \dots \circ f^{(1)}(\mathbf{x}).$$

Modularity: We can implement each layer's  
computations as a black box.




# Feature Learning

Neural nets can be viewed as a way of learning features:



The goal:




# Feature Learning

- Suppose we're trying to classify images of handwritten digits.
- Each image is represented as a vector of  $28 \times 28 = 784$  pixel values.
- Each hidden unit in the first layer acts as a **feature detector**.
- We can visualize  $\mathbf{w}$  by reshaping it into an image.  
Below is an example that responds to a diagonal stroke.



# Features for Classifying Handwritten Digits

Features learned by the first hidden layer of a handwritten digit classifier:



Unlike hard-coded feature maps (e.g., in polynomial regression), features learned by neural networks adapt to patterns in the data.

## Expressivity of a Neural Network


---



- 1 Softmax Regression
- 2 Tracking Model Performance
- 3 Limits of Linear Classification
- 4 Introducing Neural Networks
- 5 Expressivity of a Neural Network

# Expressivity

- A hypothesis space  $\mathcal{H}$  is the set of functions that can be represented by some model.
- Consider two models  $A$  and  $B$  with hypothesis spaces  $\mathcal{H}_A, \mathcal{H}_B$ .
- If  $\mathcal{H}_B \subseteq \mathcal{H}_A$ , then  $A$  is more **expressive** than  $B$ .  
 $A$  can **represent** any function  $f$  in  $\mathcal{H}_B$ .



- Some functions (XOR) can't be represented by linear classifiers.  
Are deep networks more expressive?

# Expressive Power of Linear Networks

- Consider a linear layer: the activation function was the identity. The layer just computes an affine transformation of the input.
- Any sequence of linear layers is equivalent to a single linear layer.

$$\mathbf{y} = \underbrace{\mathbf{W}^{(3)}\mathbf{W}^{(2)}\mathbf{W}^{(1)}}_{\triangleq \mathbf{W}'} \mathbf{x}$$


- Deep linear networks can only represent linear functions
  - no more expressive than linear regression.

# Expressive Power of Non-linear Networks

- Multi-layer feed-forward neural networks with non-linear activation functions
- **Universal Function Approximators:**  
They can approximate any function arbitrarily well, i.e., for any  $f : \mathcal{X} \rightarrow \mathcal{T}$  there is a sequence  $f_i \in \mathcal{H}$  with  $f_i \rightarrow f$ .
- True for various activation functions (e.g. thresholds, logistic, ReLU, etc.)

# Designing a Network to Classify XOR

Assume a hard threshold activation function.



# Designing a Network to Classify XOR

$h_1$  computes  $x_1 \vee x_2$

$$\mathbb{I}[x_1 + x_2 - 0.5 > 0]$$


$h_2$  computes  $x_1 \wedge x_2$

$$\mathbb{I}[x_1 + x_2 - 1.5 > 0]$$

$y$  computes  $h_1 \wedge (\neg h_2) = x_1 \oplus x_2$

$$\mathbb{I}[h_1 - h_2 - 0.5 > 0]$$


$$\equiv \mathbb{I}[h_1 + (1 - h_2) - 1.5 > 0]$$



# Universality for Binary Inputs and Targets

- Hard threshold hidden units, linear output
- Strategy:  $2^D$  hidden units, each of which responds to one particular input configuration


$x_1$	$x_2$	$x_3$	$t$
	$\vdots$		$\vdots$
-1	-1	1	-1
-1	1	-1	1
-1	1	1	1
	$\vdots$		$\vdots$




- Only requires one hidden layer, though it is extremely wide.

# Expressivity of the Logistic Activation Function

- What about the logistic activation function?
- Approximate a hard threshold by scaling up  $w$  and  $b$ .



$$y = \sigma(x)$$



$$y = \sigma(5x)$$


- Logistic units are differentiable, so we can learn weights with gradient descent.



# What is Expressivity Good For?

- May need a very large network to represent a function.
- Non-trivial to learn the weights that represent a function.
- If you can learn any function, over-fitting is potentially a serious concern!


For the polynomial feature mappings, expressivity increases with the degree  $M$ , eventually allowing multiple perfect fits to the training data. This motivated  $L^2$  regularization.



- Do neural networks over-fit and how can we regularize them?

# Regularization and Over-fitting for Neural Networks

- The topic of over-fitting (when & how it happens, how to regularize, etc.) for neural networks is not well-understood, even by researchers!
  - ▶ In principle, you can always apply  $L^2$  regularization.
  - ▶ You will learn more in CSC413.
- A common approach is **early stopping**, or stopping training early, because over-fitting typically increases as training progresses.



- Don't add an explicit  $\mathcal{R}(\theta)$  term to our cost.

# Conclusion

- Multi-class classification
- Selecting good metrics to track performance in models
- From linear to non-linear models